


Introduction to PyTorch

L. Palzer

Machine Learning for Communications – TUM LNT

November 12, 2019



Overview

Deep Learning Frameworks

PyTorch Tensors

PyTorch Automatic Differentiation

PyTorch Training a Neural Networks

Outlook

Deep Learning Frameworks

Why use a deep learning framework?

- Implementing own models from scratch does not scale
- Huge open source community is there to help

How to choose one?

- Ease of use
- Computation speed
- Active community
- Applications (images, NLP, etc.)

Deep Learning Frameworks

Most prominent **deep learning** frameworks for python:

- TensorFlow
- PyTorch



PYTORCH

Other frameworks

- Microsoft Cognitive Toolkit
- Caffe2 (merged into Pytorch in March 2018)
- Theano (discontinued by original developers in May 2018)
- ...

TensorFlow

- Starting in 2011, Google Brain built **DistBelief** as a proprietary machine learning framework
- TensorFlow emerged from this as an **open source library**, first release in November 2015
- Current stable release: **2.0.0** (October 1, 2019)
- **Native Python API**, but with bindings to many other languages
- Tensorflow Lite: release for **embedded devices** such as Android based smartphones



PyTorch

- Based on the **Torch**¹ library for **Lua**, first converted to Python by Adam Paszke
- Initial release in October 2016
- Now primarily developed by **Facebook AI** team
- Current stable release: **1.3.0** (October 10, 2019)
- PyTorch Mobile: experimental release for iOS and Android

PYTORCH

¹<http://torch.ch>

<https://pytorch.org>

TensorFlow vs. Pytorch

Before 2.0 / 1.3

TensorFlow:

- + Better performance
- + Support for embedded devices
- + Larger community
- + Visualization with Tensorboard
- Difficult to learn/ adapt models/ debug

→ Production

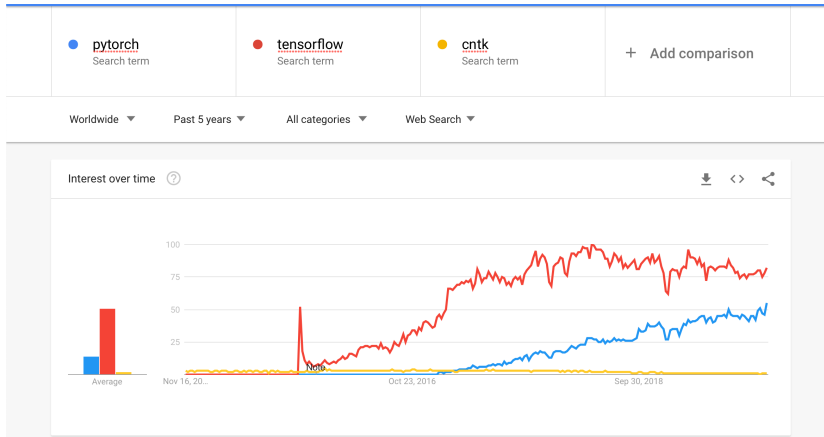
Pytorch:

- + More "Pythonic"
- + More flexible
- + Easier to debug
- + Easier to parallelize
- No proper visualization


→ Research

With 2.0 / 1.3, increasingly a "matter of taste"

Google Trend Search



Popularity in Academia



High Level APIs

Keras²:

- Developed by François Chollet at Google
- Works on top of TensorFlow, CNTK, Theano
- Recommended starting point for learning TensorFlow



fastai³⁴

- Works on top of PyTorch
- Project lead by Jeremy Howard
- Used in fast.ai online courses



²<https://keras.io>

³<https://docs.fast.ai>

⁴<https://www.fast.ai>

Example with fastai

Train a world class model with five lines of code:

```
from fastai.vision import *  
path = untar_data(MNIST_PATH)  
data = image_data_from_folder(path)  
learn = cnn_learner(data, models.resnet18, metrics=accuracy)  
learn.fit(5)
```

PyTorch Basics

https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html

Getting PyTorch

- Install PyTorch locally:
`pip3 install torch torchvision`
- Already **available in Google Colab**
- Import module:

```
>>> import torch
```

Tensors

All of you have encountered special classes of tensors before:

- **Scalars**: Zero order tensor
- **Vectors**: First order tensor
- **Matrices**: Second order tensor
- \vdots
- **Multidimensional arrays**: higher order tensor

Tensors in PyTorch are similar to NumPy's ndarrays, but tensors can be used on a GPU to accelerate computing

Basics: Tensors

- Create a tensor of zeros:

```
>>> a = torch.zeros(2,3)
>>> print(a)
tensor([[0 0 0],
        [0 0 0]])
```

- Create a tensor from data:

```
>>> b = torch.tensor([1.25,4.3])
>>> print(b)
tensor([1.25,4.3])
```

- Important methods such as `.empty`, `.randn`, `.float` etc.

Tensors from NumPy

- Pass by reference

```
>>> a = np.array([[1,2],[3,4]])
>>> print(a)
[[1 2]
 [3 4]]
>>> b = torch.from_numpy(a[:,0])
>>> print(b)
tensor([1, 3])
>> b[:] = 8
>>> print(a)
[[8 2]
 [8 4]]
```


Tensors from NumPy

- Copy from NumPy:

```
>>> a = np.array([[1,2],[3,4]])
>>> b = torch.tensor(a)
>>> c = torch.tensor(a, dtype=torch.float32) #for float
>>> print(c.numpy())
[[1 2]
 [3 4]]
```

- Indexing works the same as in NumPy
- For reshaping, use `.view` on a tensor similar to `.reshape` in NumPy

PyTorch Automatic Differentiation

https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html

Autograd

- Central package in PyTorch: `autograd`
- Provides **automatic differentiation** for all operations on tensors
- Workflow:
 - Set tensor attribute `.requires_grad` to `True` to start tracking operations
 - Perform computations as usual
 - Each involved tensor has now a function `.grad_fn` referencing the gradient
 - **Backprop** the gradients using `.backward()`
 - Call the gradient via `.grad`

Autograd Example

- Consider $y = (2 \cdot x_1 + x_2)^2$ with $x_1 = 3$ and $x_2 = 4$

```
>>> x1 = torch.tensor(3.0,requires_grad=True)
>>> x2 = torch.tensor(4.0,requires_grad=True)
>>> y = (2*x1 + x2)
>>> print(y)
tensor(10., grad_fn=<AddBackward0>)
>>> y = y**2
tensor(100., grad_fn=<PowBackward0>)
>>> y.backward()
>>> print(x1.grad,x2.grad)
tensor(40.) tensor(20.)
```

PyTorch Neural Networks

https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html

Workflow

- Define neural network
- Set up data set
- Choose loss function and optimization algorithm
- Train!

Defining the Network

- Network is child class inheriting predefined properties and methods
- Define **learnable parameters** and **forward propagation**

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # define learnable parameters
    def forward(self, x):
        # define forward propagation
        return x
```

Learnable Parameters

- Example: network with two hidden layers

```
def __init__(self):  
    super(Net, self).__init__()  
    self.h1 = nn.Linear(1,8, bias=True)  
    self.h2 = nn.Linear(8,8, bias=True)  
    self.out = nn.Linear(8,1, bias=True)
```

- Weights and biases are initialized with uniform distribution
- Many other layers⁵, e.g., `nn.conv1d`, `nn.conv2d`, etc.

⁵<https://pytorch.org/docs/stable/nn.html#convolution-layers>

Forward Propagation

- Example: network with two hidden layers

```
def forward(self, x):  
    x = F.relu(self.h1(x)) # First hidden layer  
    x = F.relu(self.h2(x)) # Second hidden layer  
    x = self.out(x) # Output layer  
    return x
```

- Can perform many other operations such as reshaping

Example

```
>>> net = Net() # Initialize network
>>> print(net)
Net(
  (h1): Linear(in_features=1, out_features=8, bias=True)
  (h2): Linear(in_features=8, out_features=8, bias=True)
  (out): Linear(in_features=8, out_features=1, bias=True)
)
>>> params = list(net.parameters())
>>> print(params[0].T) # weights in the first layer
tensor([[ -0.3591,  0.7028, -0.0415, -0.7531,  0.3992, -0.6650,
          -0.4175,  0.5087]], grad_fn=<PermuteBackward>)
```

Dataset

- Create **dataloader** object to manage dataset

```
dataset = torch.utils.data.TensorDataset(x_tensor, y_tensor)
dataloader = torch.utils.data.DataLoader(dataset,
    batch_size = 128)
```

- Example: download **MNIST** dataset

```
train_set = torchvision.datasets.MNIST('./files/',
    train=True, download=True)
train_loader = torch.utils.data.DataLoader(train_set,
    batch_size=128, shuffle=True)
```

Optimization

- Define **loss function** and **optimization algorithm**
- Examples `.BCEloss`, `.CrossEntropyloss`

```
criterion = nn.MSELoss()
```

- Choose **optimization algorithm**⁶
- Can specify learning rate, momentum, weight decay, etc.

```
optimizer = torch.optim.SGD(net.parameters(), lr=0.001)
```

⁶See this paper: <https://arxiv.org/pdf/1609.04747.pdf>

Learning

- Run the learning procedure

```
for epoch in range(num_epoch)
    for idx, (inputs, labels) in enumerate(train_loader):
        # reset optimizer
        optimizer.zero_grad()
        # Forward pass
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        # Backward and optimize
        loss.backward()
        optimizer.step()
```

Outlook

Outlook

- Using GPUs
- Other networks, e.g., convolutional
- Different optimization algorithms